

Arc in the Linux Kernel

Alice Ryhl

Rust in the Linux Kernel

- The Rust for Linux project started in 2020.
- First RFC to add Rust to Linux in April 2021.
- Actually merged in October 2022.
- I sent the Rust Binder RFC in November 2023.



What does it take for Rust to succeed in the kernel?

- Real Rust drivers used in the real world.
- Good first impressions for people coming from C.
- Do not require a specific version of rustc.
- Compile Rust with GCC.



What does it take for Rust to succeed in the kernel?

- **Real Rust drivers used in the real world.**
- Good first impressions for people coming from C.
- Do not require a specific version of rustc.
- Compile Rust with GCC.

My previous talks



What does it take for Rust to succeed in the kernel?

- Real Rust drivers used in the real world.
- **Good first impressions for people coming from C.**
- **Do not require a specific version of rustc.**
- Compile Rust with GCC.

← This talk




Goals for this talk

- The Linux Kernel needs unstable Rust features.
- This talk aims to explain why.
- Deep dive on unstable features related to custom Arc:
 - Arbitrary self types.
 - `#[derive(SmartPointer)]`.



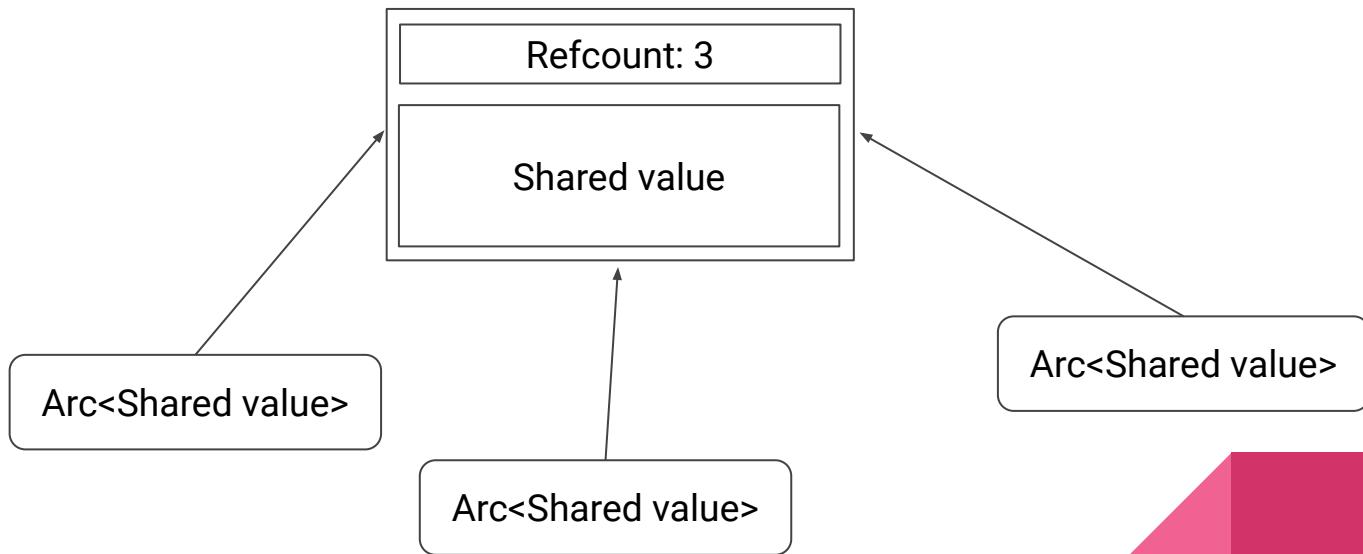
What is Arc?

- Reference counted container for sharing immutable values.




What is Arc?

- Reference counted container for sharing immutable values.




What is Arc?

- Reference counted container for sharing immutable values.



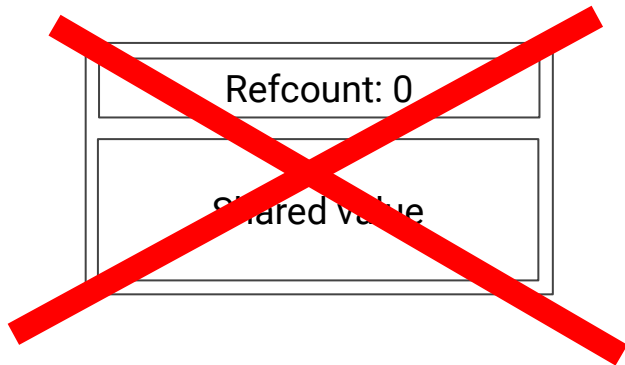
What is Arc?

- Reference counted container for sharing immutable values.



What is Arc?

- Reference counted container for sharing immutable values.



Linux has a custom Arc

- Linux can't use the standard library Arc.
- But the standard library Arc is special.
- You can't implement your own Arc.
- Linux uses unstable features to work around this.



Do not call abort on overflow

- When you call abort in the kernel, your device turns off.
 - Even if something goes wrong, it's better to try and carry on.
- When the refcount hits `isize::MAX`, replace it with ∞ .
 - (and print a warning)
- Once it hits ∞ , the refcount stays there forever.



The Linux Kernel Memory Model

- Arc uses an atomic integer for the refcount.
- Rust uses the C++ memory model for atomics.
- Linux uses the Linux Kernel Memory Model for atomics.
- Rust in Linux must also use LKMM for atomics.
 - We use inline assembly for all atomic operations.



Nice to have: Pin every Arc

- Pretty much every C type is !Unpin.
 - This includes Mutex.
- Linux's custom Arc will pin all values automatically.



Nice to have: No weak pointers

- The standard library Arc has weak references.
- Not needed in the Linux Kernel.
- Removing them simplifies our Arc.



Other custom types similar to Arc

Types that are similar to Arc. They need the same unstable features.

- ArcBorrow - similar to `&Arc<T>`
- UniqueArc - Arc with mutable access
- ListArc - Version of Arc for linked lists



Arbitrary self types

```
impl MyStruct {  
    fn my_func1(&self) { ... }  
    fn my_func2(&mut self) { ... }  
  
    fn my_func3(self: Arc<Self>) {  
        ...  
    }  
}
```

This is unstable.



Arbitrary self types

```
impl MyStruct {  
    fn my_func1(&self) { ... }  
    fn my_func2(&mut self) { ... }  
  
    fn my_func3(self: Arc<Self>) {  
        let arc = self.clone();  
    }  
}
```

Makes this possible.



Linked lists also pose a problem

- Linked lists are:
 - Hard to implement in Rust.
 - And perform worse than Vec.
- Yet, the Linux Kernel still uses them (with Arc!)



Atomic context

```
let guard = spin_lock.lock();  
guard.list.insert(value);  
drop(guard);
```

While spinlock is locked, you
can't allocate memory.

How do you implement
insert without allocating
memory?



With a vector

- If there is no more space:
 - a. Exit the spinlock.
 - b. Allocate memory for a larger vector.
 - c. Re-enter the spinlock.
 - d. Move all elements to larger vector.

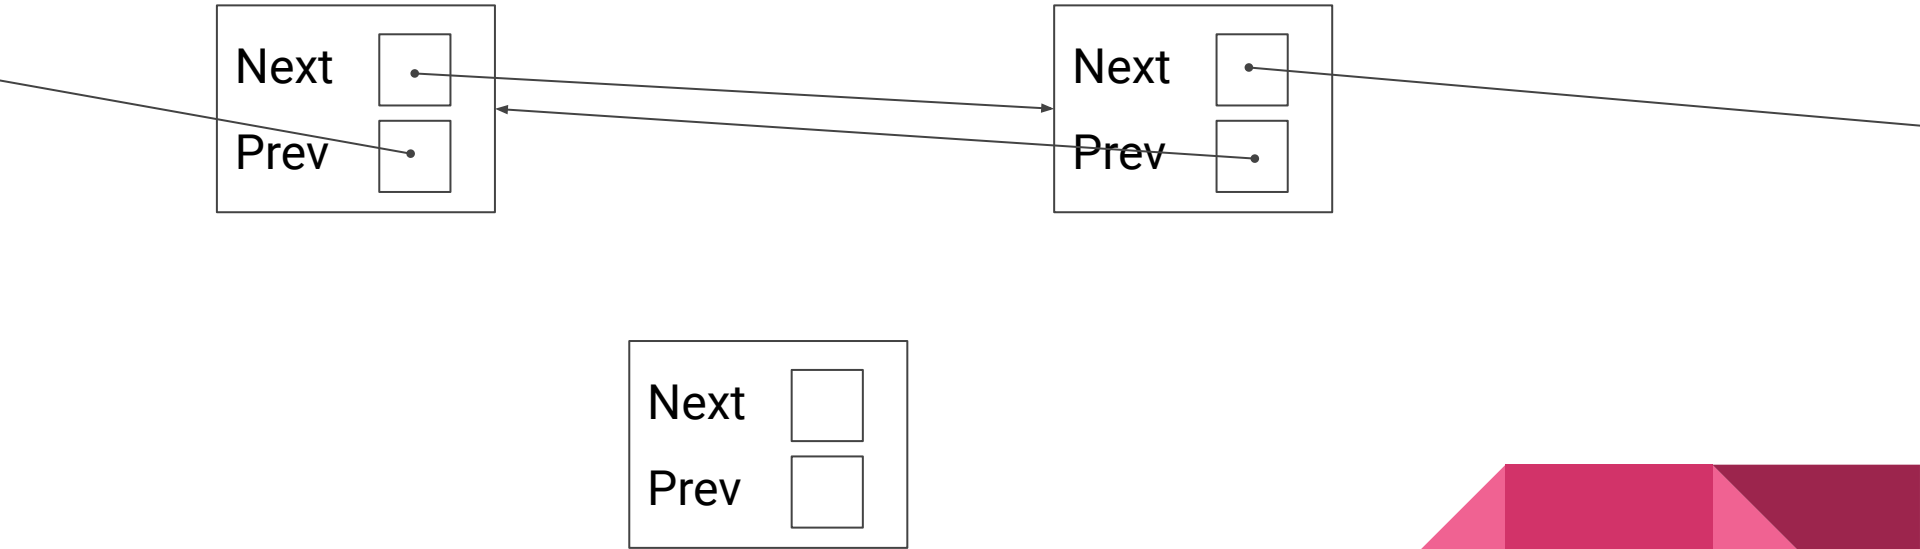


With a linked list

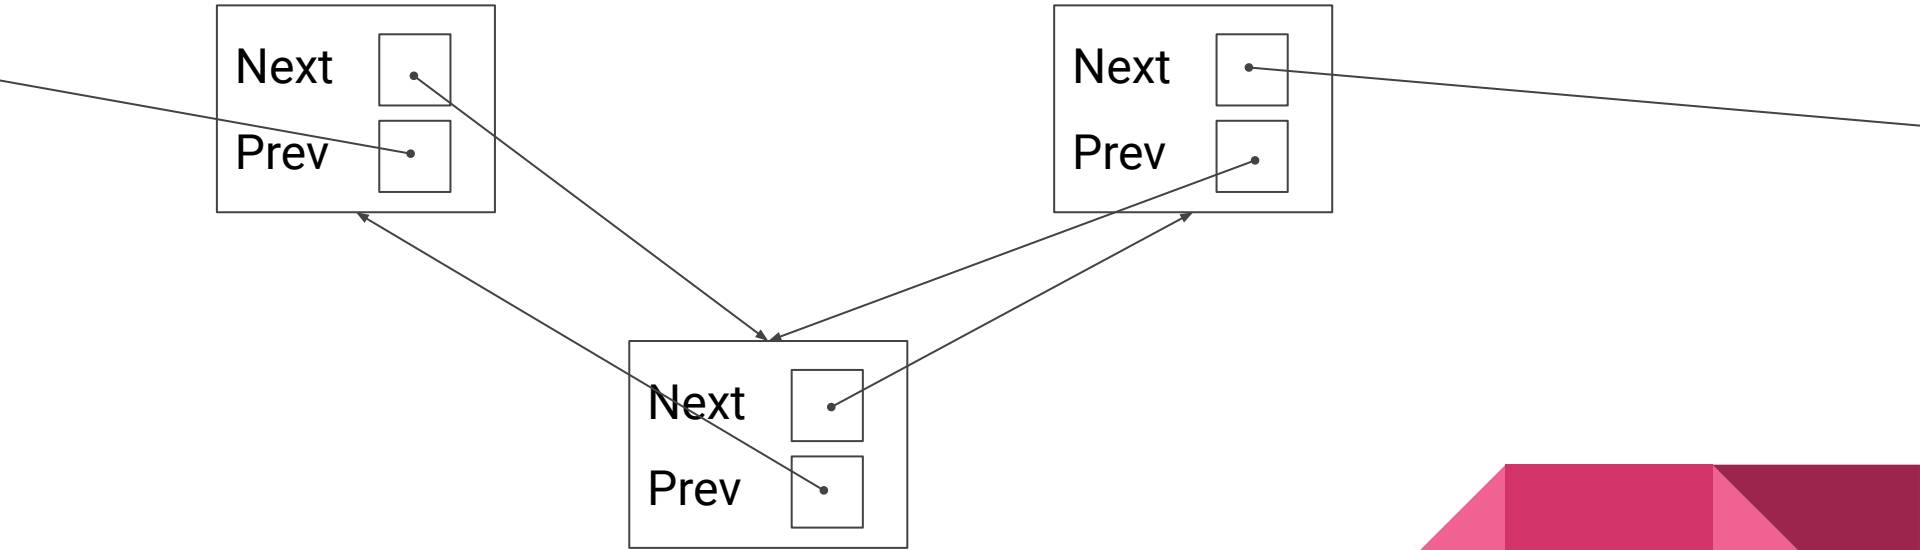
- To insert, modify prev/next pointers.
- That's it!
- No allocations needed.



Linked list



Linked list



Linked list

```
struct MyValue {
```

```
    // In practice, these are wrapped into one field using a struct.
```

```
    next: *mut MyValue,
```

```
    prev: *mut MyValue,
```

```
    foo: Foo,
```

```
    bar: Bar,
```

```
}
```

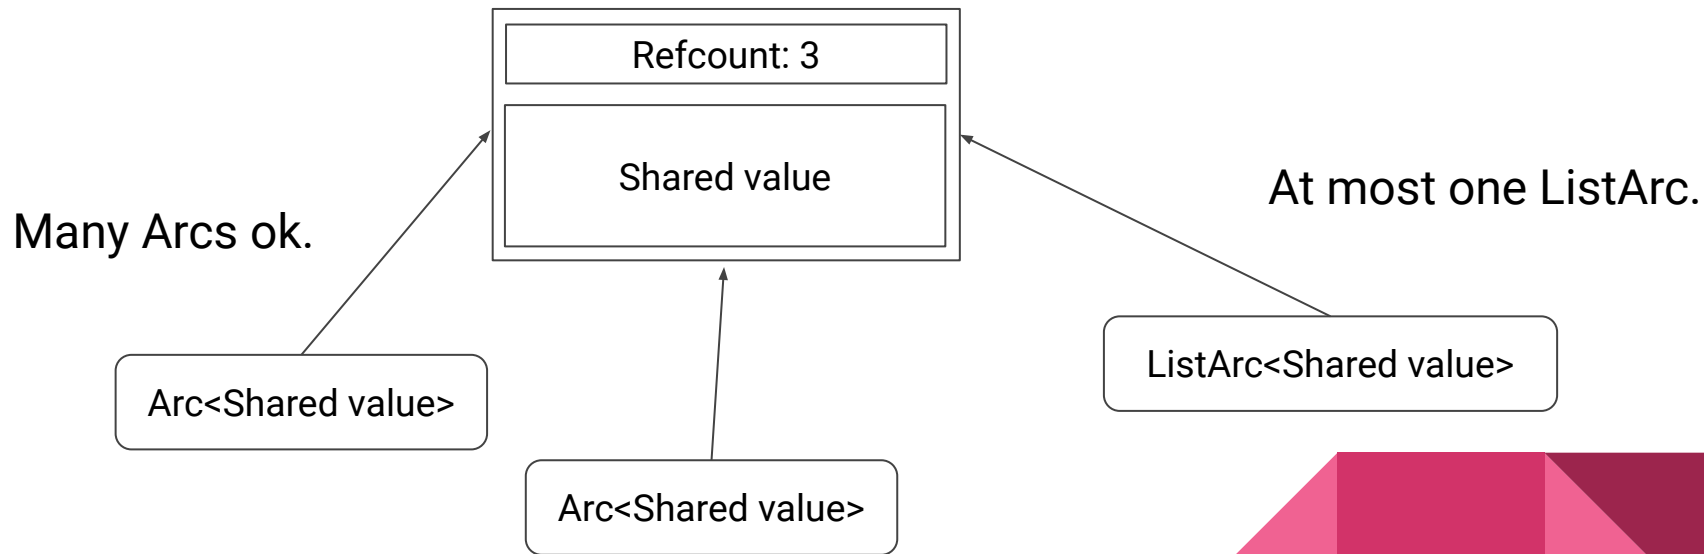


Using Arc for linked list elements

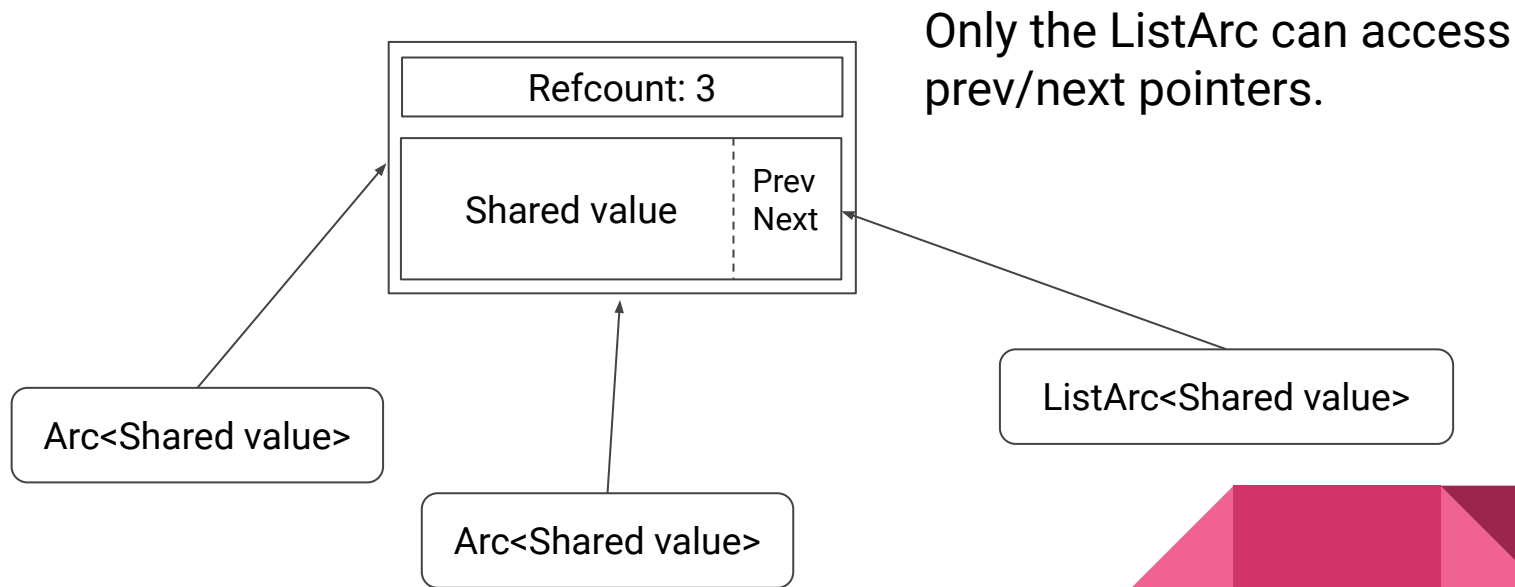
- There is a problem:
- You can have multiple Arcs to the same shared value.
- But you only have one pair of prev/next pointers.
- What if you insert into two different lists in parallel?



Solution: the ListArc



Solution: the ListArc



Queue of events

- Linked lists are used for queues of events.
- Many different types of events.
 - `LinkedList<dyn EventTrait>`



Trait for events

```
trait EventTrait {  
    fn run_event(self: ListArc<Self>);  
}
```

Can't just be &self.



Traits with custom Arc

- Using `dyn Trait` with a custom `Arc` is also unstable.
- Requires many unstable features.



Traits with Arc

```
impl<T, U> CoerceUnsize<Arc<U>> for Arc<T>
```

where

```
    T: Unsize<U>,
```

```
    T: ?Sized,
```

```
    U: ?Sized,
```

```
{}
```

Unstable



Traits with Arc

```
impl<T, U> DispatchFromDyn<Arc<U>> for Arc<T>
```

where

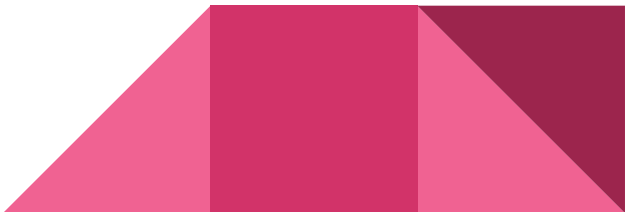
```
    T: Unsize<U>,
```

```
    T: ?Sized,
```

```
    U: ?Sized,
```

```
{}
```

Unstable



Traits with Arc

```
impl<T, U> DispatchFromDyn<Arc<U>> for Arc<T>
```

where

```
    T: Unsize<U>,
```

```
    T: ?Sized,
```

```
    U: ?Sized,
```

Unstable

```
{}
```



Solution: `#[derive(SmartPointer)]`

- We stabilize a derive macro.
- We do *not* stabilize what it expands to.
- This allows us to change the underlying traits in the future.



RFC: `#[derive(SmartPointer)]` #3621



Open

Darksonn wants to merge 1 commit into `rust-lang:master` from `Darksonn:derive-smart-pointer`



Conversation 1



Commits 1



Checks 0



Files changed 1



Darksonn posted just before RustNL • edited ▾



Use custom smart pointers with trait objects.

[Rendered](#)

Tracking issue: [rust-lang/rust#123430](#)

Co-authored by [@Darksonn](#) and [@Veykril](#)
Thank you to [@compiler-errors](#) for [the original idea](#)



RFC: `#[derive(SmartPointer)]` ...

061c786