

Compression Carcinized

implementing zlib-rs


Folkert de Vries, RustNL 2024

carcinization



a form of convergent evolution in which
non-crab crustaceans evolve a **crab-like**
body plan

zlib: a library you've used today



zlib: a library you've used today

×	Headers	Preview	Response	Initiator	Timing	Cookies
▼	General					
Request URL:	https://www.rust-lang.org/					
Request Method:	GET					
Status Code:	● 200 OK					
Remote Address:	18.239.94.125:443					
Referrer Policy:	origin					
▼	Response Headers					
Content-Encoding:	gzip					
Content-Security-Policy:	default-src 'self'; frame-ancestors					
Content-Type:	text/html; charset=utf-8					
Date:	Thu, 11 Apr 2024 10:23:13 GMT					

zlib: a library you've used today

```
> objdump -T /usr/lib/x86_64-linux-gnu/libz.so | grep "compress"
00000000000010370 g      DF .text 00000000000000022 ZLIB_1.2.0 compressBound
00000000000010360 g      DF .text 0000000000000000f Base      compress
00000000000010220 g      DF .text 00000000000000013d Base      compress2
00000000000010560 g      DF .text 00000000000000001c Base      uncompress
000000000000103a0 g      DF .text 00000000000000001c0 ZLIB_1.2.9 uncompress2
```

```
pub unsafe extern "C" fn compress2(
    dest: *mut Bytef,
    destLen: *mut c_ulong,
    source: *const Bytef,
    sourceLen: c_ulong,
    level: c_int
) -> c_int
```

project goals



drop-in replacement for
the zlib dynamic library



high-performance
implementation for rust

Topics

- crash course compression
- the zlib ecosystem
- ponderings on porting

crash course compression

Why use **many** byte
when **few** do trick?

Why Compress?

cost



speed



Lossless Compression

```
assert_eq!(decompress(compress(data)), data)
```

Recognizing patterns

foobarfoo



foobar<offset = 6, len = 3>

Finding patterns

3.141592653589793238462643383279502884197169399
3751058209749445923078164062862089986280348253
4211706798214808651328230664709384460955058223
176

goal: find the (longest)
<offset,len> insertions

Finding patterns

s	e	r	i	e	u	s	p	r	o
---	---	---	---	---	---	---	---	---	---




Finding patterns

s	e	r	i	e	u	s	p	r	o
---	---	---	---	---	---	---	---	---	---




Finding patterns

The **window size** determines how far back the offset can go



Finding patterns

The **compression level** determines how hard we try to find the longest match



Finding patterns

f	o	o	b	a	r	f	o	o	...
---	---	---	---	---	---	---	---	---	-----



Finding patterns

"foo" \rightarrow { 0 }


f	o	o	b	a	r	f	o	o	...
---	---	---	---	---	---	---	---	---	-----



Finding patterns

"foo" \rightarrow { 0 }

"oob" \rightarrow { 1 }



Finding patterns

"foo" -> { 0 }

"oob" -> { 1 }

"oba" -> { 2 }

"bar" -> { 3 }

"arf" -> { 4 }

"rfo" -> { 5 }




Finding patterns



very effective for web data, even at low
compression levels

Streaming



zlib can stream compression and
decompression

zlib & rust

what are we even implementing

zlib-adler: the OG



goal: **stability**

still supports 16-bit
systems

does not use modern
hardware well

zlib-ng: the next generation



goal: **performance**

removes legacy,
but API-compatible

uses SIMD to speed up
the algorithm

why rust



reduced surface area

why rust

*“ Any sufficiently complicated C project
contains an
ad hoc,
informally-specified,
bug-ridden,
slow
implementation of half of cargo*

flate2

flate2 v1.0.30

DEFLATE compression and decompression exposed as Read/BufRead/Write streams. Supports miniz_oxide and multiple zlib implementations. Supports zlib, gzip, and raw deflate streams.

[Homepage](#) [Documentation](#) [Repository](#)



All-Time: 124,987,458



Recent: 19,092,826



Updated: 5 days ago

a nice rust API for zlib

used in cargo

miniz-oxide: better safe than sorry




goal: **safety**

a safe (but slow) rust
implementation

does not cover the full
zlib API

zlib-rs: a safer zlib



goal: **safety & performance**

faster through the
use of SIMD

implements the full zlib
API

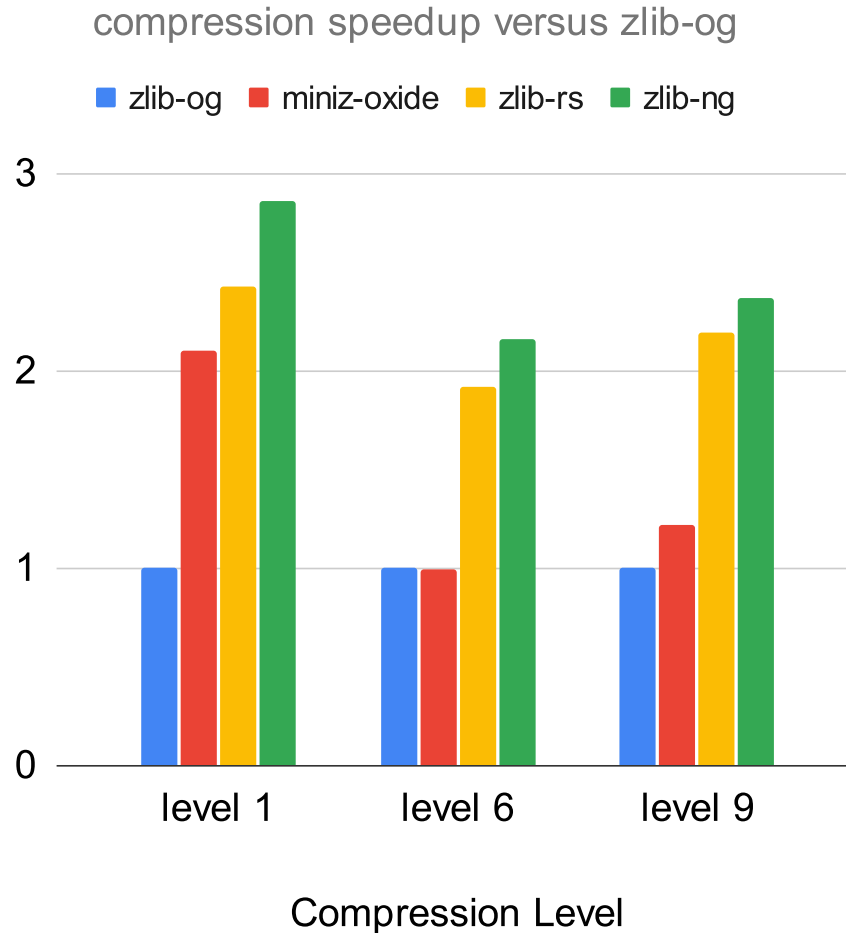
unsafe sandwich

unsafe C API

(mostly) safe business logic

unsafe SIMD

compression speed





progress

```
[dependencies]
flate2 = {
  version = "1.0.29",
  default-features = false,
  features = ["zlib-rs"]
}
```

early days, but give it a go!

a crab-like body plan



from C to rust

spectrum of porting



implementation

rewrite



spectrum of porting



implementation

rewrite



implementation

- e.g. rustls, ntpd-rs
- reinvent the wheel
- innovate on architecture
- high risk, high reward

spectrum of porting



implementation

rewrite



implementation

- e.g. rustls, ntpd-rs
- reinvent the wheel
- innovate on architecture
- high risk, high reward

Rewrite

- e.g. rav1d
- use existing knowledge
- inherits architecture
- works on day 1

spectrum of porting



implementation

rewrite

```
// ...  
i = 0 as libc::c_int;  
while i < nblock {  
    ftab[*eclass8.offset(i as isize) as usize] += 1;  
    ftab[*eclass8.offset(i as isize) as usize];  
    i += 1;  
    i;  
}  
// ...
```

spectrum of porting



implementation

rewrite



zlib

- reuse existing knowledge (correctness & performance)
- quick results
- architecture is constrained anyway

RiiR and You

compatability



just better




RiiR and You

funding



adoption





Summary

why use **many** bytes when **few** do trick

unreasonably effective on web content

try **zlib-rs**

use more (unglamorous) rust **in production**

 evolve crab-like body plans 

Thanks

github.com/memorysafety/zlib-rs

memorysafety.org/initiative/zlib/

slides.com/folkertdevries/compression-carcinized

is it bounds checks?

Benchmark 1 (42 runs): target/release/examples/compress 1 rs silesia-small.tar						
measurement	mean ± σ	min ... max	outliers	delta		
wall_time	119ms ± 1.97ms	117ms ... 128ms	1 (2%)	0%		
peak_rss	26.7MB ± 85.7KB	26.6MB ... 26.9MB	0 (0%)	0%		
cpu_cycles	406M ± 4.67M	399M ... 424M	1 (2%)	0%		
instructions	660M ± 469	660M ... 660M	0 (0%)	0%		
cache_references	8.06M ± 1.31M	5.65M ... 11.3M	0 (0%)	0%		
cache_misses	461K ± 36.5K	433K ... 555K	5 (12%)	0%		
branch_misses	3.59M ± 6.42K	3.58M ... 3.61M	1 (2%)	0%		
Benchmark 2 (43 runs): removed-bounds/release/examples/compress 1 rs silesia-small.tar						
measurement	mean ± σ	min ... max	outliers	delta		
wall_time	118ms ± 2.53ms	115ms ... 127ms	3 (7%)	- 1.3% ± 0.8%		
peak_rss	26.8MB ± 77.9KB	26.7MB ... 27.0MB	0 (0%)	+ 0.2% ± 0.1%		
cpu_cycles	400M ± 8.00M	391M ... 437M	2 (5%)	- 1.4% ± 0.7%		
instructions	623M ± 522	623M ... 623M	0 (0%)	⚡ - 5.6% ± 0.0%		
cache_references	7.91M ± 1.45M	5.89M ... 11.9M	1 (2%)	- 1.9% ± 7.4%		
cache_misses	458K ± 29.1K	433K ... 550K	1 (2%)	- 0.5% ± 3.1%		
branch misses	3.34M ± 7.35K	3.33M ... 3.36M	0 (0%)	⚡ - 6.8% ± 0.1%		

<https://blog.readysset.io/bounds-checks/>
<https://github.com/andrewrk/poop>

Lossy Compression



622kb

Lossy Compression



622kb ➡ 87kb